

Exploratory Data Analysis of the Forest Fires Data Set

Summary of the data set

The data used in the project contains information about the burnt areas of forest fires located in the northeast region of Portugal created by Cortez and Morais (2007). The data source is from the UCI Machine Learning Repository (Dua and Graff 2017) and can be accessed [here](#). There are 517 observations and 13 rows in the data, and there are no missing values in the dataset. Each row represents one fire monitoring instance, with the column area as our target (showing the burned area), and 12 other measurements and indexes as features (including month, day, RH, rain, DC, ISI etc).

For the exploratory data analysis, we try to follow the EDA checklist in chapter 4 of the [Art of Data Science](#).

1. Formulate your question

Can we predict the burned areas of forests using meteorological observations, soil moisture indices, and locational data?

2. Read in the data

```
In [1]: import pandas as pd
import numpy as np
import altair as alt
from sklearn.model_selection import train_test_split
alt.renderers.enable("mimetype")
```

```
Out[1]: RendererRegistry.enable('mimetype')
```

```
In [2]: forest_data = pd.read_csv("../data/raw/forestfires.csv")
forest_data
```

```
Out[2]:
```

	X	Y	month	day	FFMC	DMC	DC	ISI	temp	RH	wind	rain	area
0	7	5	mar	fri	86.2	26.2	94.3	5.1	8.2	51	6.7	0.0	0.00
1	7	4	oct	tue	90.6	35.4	669.1	6.7	18.0	33	0.9	0.0	0.00
2	7	4	oct	sat	90.6	43.7	686.9	6.7	14.6	33	1.3	0.0	0.00
3	8	6	mar	fri	91.7	33.3	77.5	9.0	8.3	97	4.0	0.2	0.00
4	8	6	mar	sun	89.3	51.3	102.2	9.6	11.4	99	1.8	0.0	0.00
...
512	4	3	aug	sun	81.6	56.7	665.6	1.9	27.8	32	2.7	0.0	6.44
513	2	4	aug	sun	81.6	56.7	665.6	1.9	21.9	71	5.8	0.0	54.29
514	7	4	aug	sun	81.6	56.7	665.6	1.9	21.2	70	6.7	0.0	11.16
515	1	4	aug	sat	94.4	146.0	614.7	11.3	25.6	42	4.0	0.0	0.00
516	6	3	nov	tue	79.5	3.0	106.7	1.1	11.8	31	4.5	0.0	0.00

3. Check the packaging

In this section, we check the packaging of the data to get some information about the data.

```
In [3]: # Number of columns and rows
forest_data.shape
```

```
Out[3]: (517, 13)
```

There are 517 observations and 13 columns.

```
In [4]: # Information for each column
forest_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 517 entries, 0 to 516
Data columns (total 13 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0    X           517 non-null    int64
 1    Y           517 non-null    int64
 2   month       517 non-null    object
 3    day         517 non-null    object
 4   FFMC        517 non-null    float64
 5    DMC         517 non-null    float64
 6    DC           517 non-null    float64
 7    ISI          517 non-null    float64
 8   temp        517 non-null    float64
 9    RH           517 non-null    int64
10   wind         517 non-null    float64
11   rain         517 non-null    float64
12   area         517 non-null    float64
dtypes: float64(8), int64(3), object(2)
memory usage: 52.6+ KB
```

By using `info()`, we get information about the data including the names of each column, non-null values, the data types of each column and the memory usage.

```
In [5]: # Missing values for each column
forest_data.isnull().sum()
```

```
Out[5]: X           0
        Y           0
        month       0
        day         0
        FFMC        0
        DMC         0
        DC          0
        ISI         0
        temp        0
        RH          0
        wind        0
        rain        0
        area        0
dtype: int64
```

We see that our data does not contain any missing values.

4. Looking at the top and the bottom of the data

After checking the packaging, it is useful to look at the beginning and end of the dataset. This allows us to know if the data was read in properly and if things are properly formatted.

```
In [6]: forest_data.head()
```

```
Out[6]:
```

	X	Y	month	day	FFMC	DMC	DC	ISI	temp	RH	wind	rain	area
0	7	5	mar	fri	86.2	26.2	94.3	5.1	8.2	51	6.7	0.0	0.0
1	7	4	oct	tue	90.6	35.4	669.1	6.7	18.0	33	0.9	0.0	0.0
2	7	4	oct	sat	90.6	43.7	686.9	6.7	14.6	33	1.3	0.0	0.0
3	8	6	mar	fri	91.7	33.3	77.5	9.0	8.3	97	4.0	0.2	0.0
4	8	6	mar	sun	89.3	51.3	102.2	9.6	11.4	99	1.8	0.0	0.0

```
In [7]: forest_data.tail()
```

```
Out[7]:
```

	X	Y	month	day	FFMC	DMC	DC	ISI	temp	RH	wind	rain	area
512	4	3	aug	sun	81.6	56.7	665.6	1.9	27.8	32	2.7	0.0	6.44
513	2	4	aug	sun	81.6	56.7	665.6	1.9	21.9	71	5.8	0.0	54.29
514	7	4	aug	sun	81.6	56.7	665.6	1.9	21.2	70	6.7	0.0	11.16
515	1	4	aug	sat	94.4	146.0	614.7	11.3	25.6	42	4.0	0.0	0.00
516	6	3	nov	tue	79.5	3.0	106.7	1.1	11.8	31	4.5	0.0	0.00

5. Check the "n"s

In this section, we check counts and unique values of the categorical variables and calculate some summary statistics for the numerical columns.

Before we start, we split our data into 80% training and 20% testing sets and will only perform the exploratory data analysis on the training portion of our data.

```
In [8]: # split data into training and testing
train_set, test_set = train_test_split(
    forest_data, test_size = 0.2, random_state = 123
)
```

```
In [9]: # number of observations in the training set
train_set.shape
```

```
Out[9]: (413, 13)
```

We have 413 observations and 13 columns in the training set.

We start by counting the unique values in categorical columns. We can see that the x-axis spatial coordinate within the Montesinho park (X) ranges from 1 to 9. The y-axis spatial coordinate within the Montesinho park (Y) ranges from 2 to 9, however there is no y-coordinate of 7.

```
In [10]: train_set["X"].value_counts()
```

```
Out[10]: 6      74
         4      71
         2      56
         7      49
         8      46
         3      46
         1      40
         5      24
         9       7
         Name: X, dtype: int64
```

```
In [11]: train_set["Y"].value_counts()
```

```
Out[11]: 4      164
         5      104
         6       53
         3       52
         2       37
         9        2
         8         1
         Name: Y, dtype: int64
```

Next, we count the unique values in month and day. We can see that majority of the observations are in the months of August and September. Moreover, the data has more observations during the weekends compared to the weekdays.

```
In [12]: train_set["month"].value_counts()
```

```
Out[12]: aug      149
         sep      138
         mar       45
         jul       22
         feb       17
         jun       13
         oct       12
         apr        8
         dec        6
         jan        1
         nov        1
         may        1
         Name: month, dtype: int64
```

```
In [13]: train_set["day"].value_counts() # Mostly during weekend
```

```
Out[13]: sun      73
         fri      70
         sat      69
         mon      57
         thu      51
         tue      48
         wed      45
         Name: day, dtype: int64
```

Next, we calculate some summary statistics of our data. We can see the ranges of values, as well as the means and standard deviations.

```
In [14]: # summary statistics
         train_set.describe()
```

Out [14]:	X	Y	FFMC	DMC	DC	ISI	temp	RH	
count	413.000000	413.000000	413.000000	413.000000	413.000000	413.000000	413.000000	413.000000	4
mean	4.629540	4.237288	90.771429	109.854237	546.031235	8.971671	18.819613	44.353511	
std	2.278178	1.164551	4.655424	63.576254	251.835608	4.581362	5.789594	16.476107	
min	1.000000	2.000000	50.400000	2.400000	7.900000	0.400000	4.200000	15.000000	
25%	3.000000	4.000000	90.200000	61.100000	433.300000	6.400000	15.400000	33.000000	
50%	4.000000	4.000000	91.700000	108.000000	664.500000	8.400000	19.300000	42.000000	
75%	6.000000	5.000000	92.900000	141.300000	713.900000	10.700000	22.900000	53.000000	
max	9.000000	9.000000	96.200000	291.300000	860.600000	56.100000	33.300000	99.000000	

Here, we calculate the correlation matrix for the numerical columns of our data. From the summary statistics above, we can see that many of the columns contain outliers, and thus we do not use the default method of the `corr()` function (Pearson's correlation coefficient), as it is very sensitive to outliers. We use the Spearman's rank correlation coefficient, as it is more robust to the effect of outliers.

```
In [15]: # Correlation matrix
train_set.corr(method = 'spearman')
```

Out [15]:		X	Y	FFMC	DMC	DC	ISI	temp	RH	wind
X	1.000000	0.462903	-0.036634	-0.063781	-0.062645	0.000960	-0.091185	0.110238	0.054414	
Y	0.462903	1.000000	0.058250	0.017757	-0.085801	0.034649	-0.030783	0.040472	-0.001240	
FFMC	-0.036634	0.058250	1.000000	0.523017	0.257547	0.776041	0.613044	-0.330960	-0.049079	
DMC	-0.063781	0.017757	0.523017	1.000000	0.566408	0.429883	0.535974	0.017812	-0.121018	
DC	-0.062645	-0.085801	0.257547	0.566408	1.000000	0.094943	0.316728	0.041631	-0.231783	
ISI	0.000960	0.034649	0.776041	0.429883	0.094943	1.000000	0.426436	-0.190883	0.113262	
temp	-0.091185	-0.030783	0.613044	0.535974	0.316728	0.426436	1.000000	-0.496152	-0.164380	
RH	0.110238	0.040472	-0.330960	0.017812	0.041631	-0.190883	-0.496152	1.000000	0.011832	
wind	0.054414	-0.001240	-0.049079	-0.121018	-0.231783	0.113262	-0.164380	0.011832	1.000000	
rain	0.115815	0.072720	0.121794	0.121047	0.023028	0.122755	0.054985	0.182781	0.110012	
area	0.116381	0.052622	0.058037	0.082359	0.074541	0.051099	0.091627	-0.050826	0.052180	

6. Making plots

In this section, we make plots to further our understanding of data.

Dependent Variable (area)

We start by plotting our target variable (the burnt area of the forest).

```
In [16]: alt.Chart(train_set).mark_bar().encode(
    alt.X("area", bin = alt.Bin(maxbins = 30)),
    y = "count() "
)
```

Out[16]:




Fig. 1 Histogram of the burnt area of the forest

From figure 1, we see that the target variable is highly skewed, and there are many observations with burnt area of 0. As a result, we apply log transformation to the target variable.

In [17]:

```
# log transformation
train_set["log_area"] = np.log(train_set["area"].replace(0, np.nan))

alt.Chart(train_set).mark_bar().encode(
    alt.X("log_area",
        bin = alt.Bin(maxbins = 20),
        title = "Area Burnt (Log Transformation)",
        y = alt.Y("count()")
    )
)
```

Out[17]:

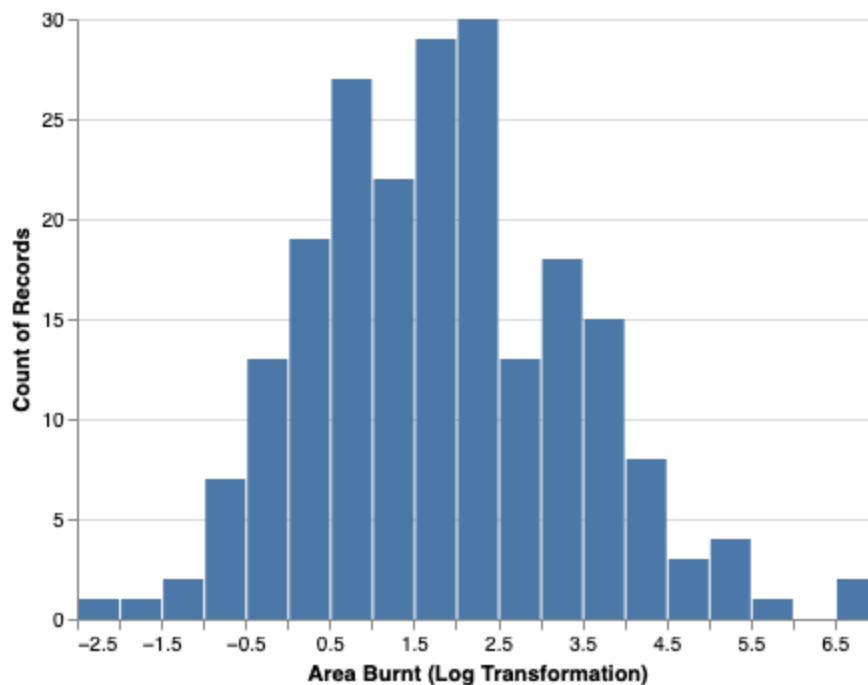


Fig 2. Histogram of the burnt area of the forest (log transformed)

From figure 2, we see that after log transforming the target variable, the target variable looks approximately normal.

Predictors

After examining the target variable, we explore our categorical and numerical variables.

Categorical Columns

As discussed earlier, many of the observations have a burnt area of 0, and consequently for visualization purposes, we apply a square root transformation to the area burnt for the visualization of the categorical variables.

In [18]:

```
alt.Chart(train_set).mark_boxplot(size = 15).encode(
    x = alt.X("area",
              scale = alt.Scale(type = "sqrt"),
              title = "Area Burnt (Square Root Transformation)"),
    y = alt.Y("day",
              sort = "x",
              title = "Day of Week"),
    color = alt.Color("day",
                      legend = None)
).properties(
    height = 250,
    width = 450
)
```

Out[18]:

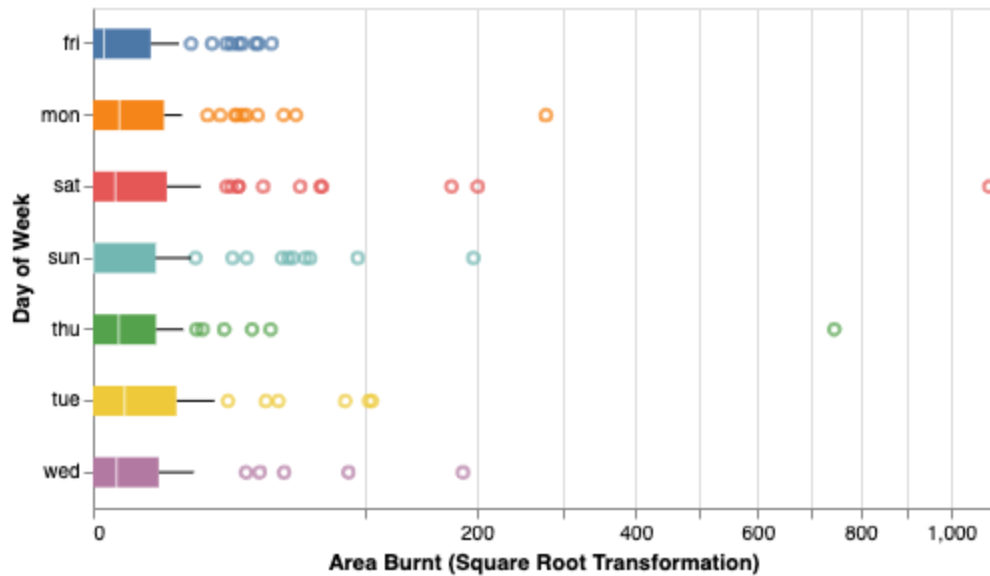


Fig 3. Boxplots of burnt areas of the forest (sqrt transformed) per day of the week

Figure 3 shows that there is no clear relationship between the burnt area of the forest and the days of the week.

In [19]:

```
alt.Chart(train_set).mark_boxplot(size = 15).encode(
    x = alt.X("area",
              scale = alt.Scale(type = "sqrt"),
              title = "Area Burnt (Square Root Transformation)"),
    y = alt.Y("month",
              sort = "x",
              title = "Month"),
    color = alt.Color("month",
                      legend = None)
).properties(
    height = 250,
    width = 450
)
```

Out[19]:

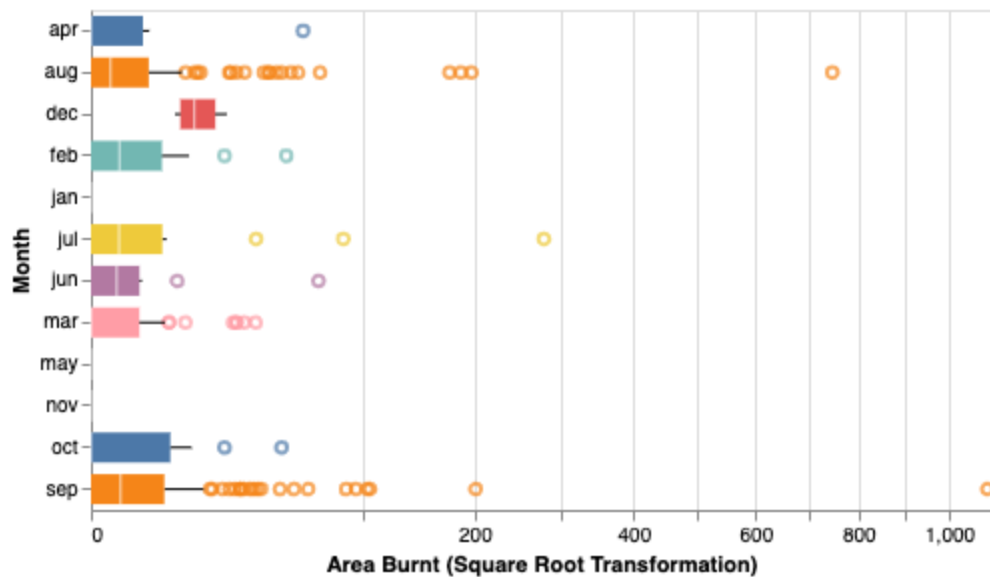


Fig 4. Boxplots of burnt areas of the forest (sqrt transformed) per month

Figure 4 shows that some months such as January, May and November do not have many observations, and since the month variable is unbalanced, to avoid overfitting, we create a `season` variable.

In [20]:

```
season_mapping = {
    "dec" : "winter",
    "jan" : "winter",
    "feb" : "winter",
    "mar" : "spring",
    "apr" : "spring",
    "may" : "spring",
    "jun" : "summer",
    "jul" : "summer",
    "aug" : "summer",
    "sep" : "fall",
    "oct" : "fall",
    "nov" : "fall"
}
train_set["season"] = train_set["month"].map(season_mapping)

alt.Chart(train_set).mark_boxplot(size = 20).encode(
    x = alt.X("area",
        scale = alt.Scale(type = "sqrt"),
        title = "Area Burnt (Square Root Transformation)"),
    y = alt.Y("season",
        sort = "x",
        title = "Season"),
    color = alt.Color("season",
        legend = None)
).properties(
    height = 200,
    width = 450
)
```

Out[20]:




Fig 5. Boxplots of burnt areas of the forest (sqrt transformed) per season

From figure 5, we can see that there is a distinction between the burnt areas of forest and seasons.

Numeric Columns

In [21]:

```
alt.Chart(train_set).mark_circle().encode(
    x = alt.X("X",
              title = "X-axis Spacial Coordinate"),
    y = alt.Y("Y",
              title = "Y-axis Spacial Coordinate"),
    size = alt.Size("area",
                   scale = alt.Scale(range = (20, 1500)),
                   title = "Burnt Area")
).configure_mark(
    color = "red",
    opacity = 0.7
)
```

Out [21]:




Fig 6. Location of the burnt areas of the forest

Figure 6 shows the location and size of the burnt areas in the forest. We see the different locations of the park have experienced different extents of areas burnt. We can see that some areas of the forest have definitely experienced larger fires. In particular coordinates of (6, 5) and (8, 6) stand out.

```

In [22]: alt.Chart(train_set).mark_circle().encode(
    x = alt.X(alt.repeat("row"), type = "quantitative"),
    y = alt.Y(alt.repeat("column"), type = "quantitative"),
    color = "season"
).properties(
    width = 110,
    height = 110
).repeat(
    column = ["FFMC", "DMC", "DC", "ISI", "temp", "RH", "wind", "rain"],
    row = ["FFMC", "DMC", "DC", "ISI", "temp", "RH", "wind", "rain"]
).configure_mark(
    opacity = 0.4
).interactive()

```

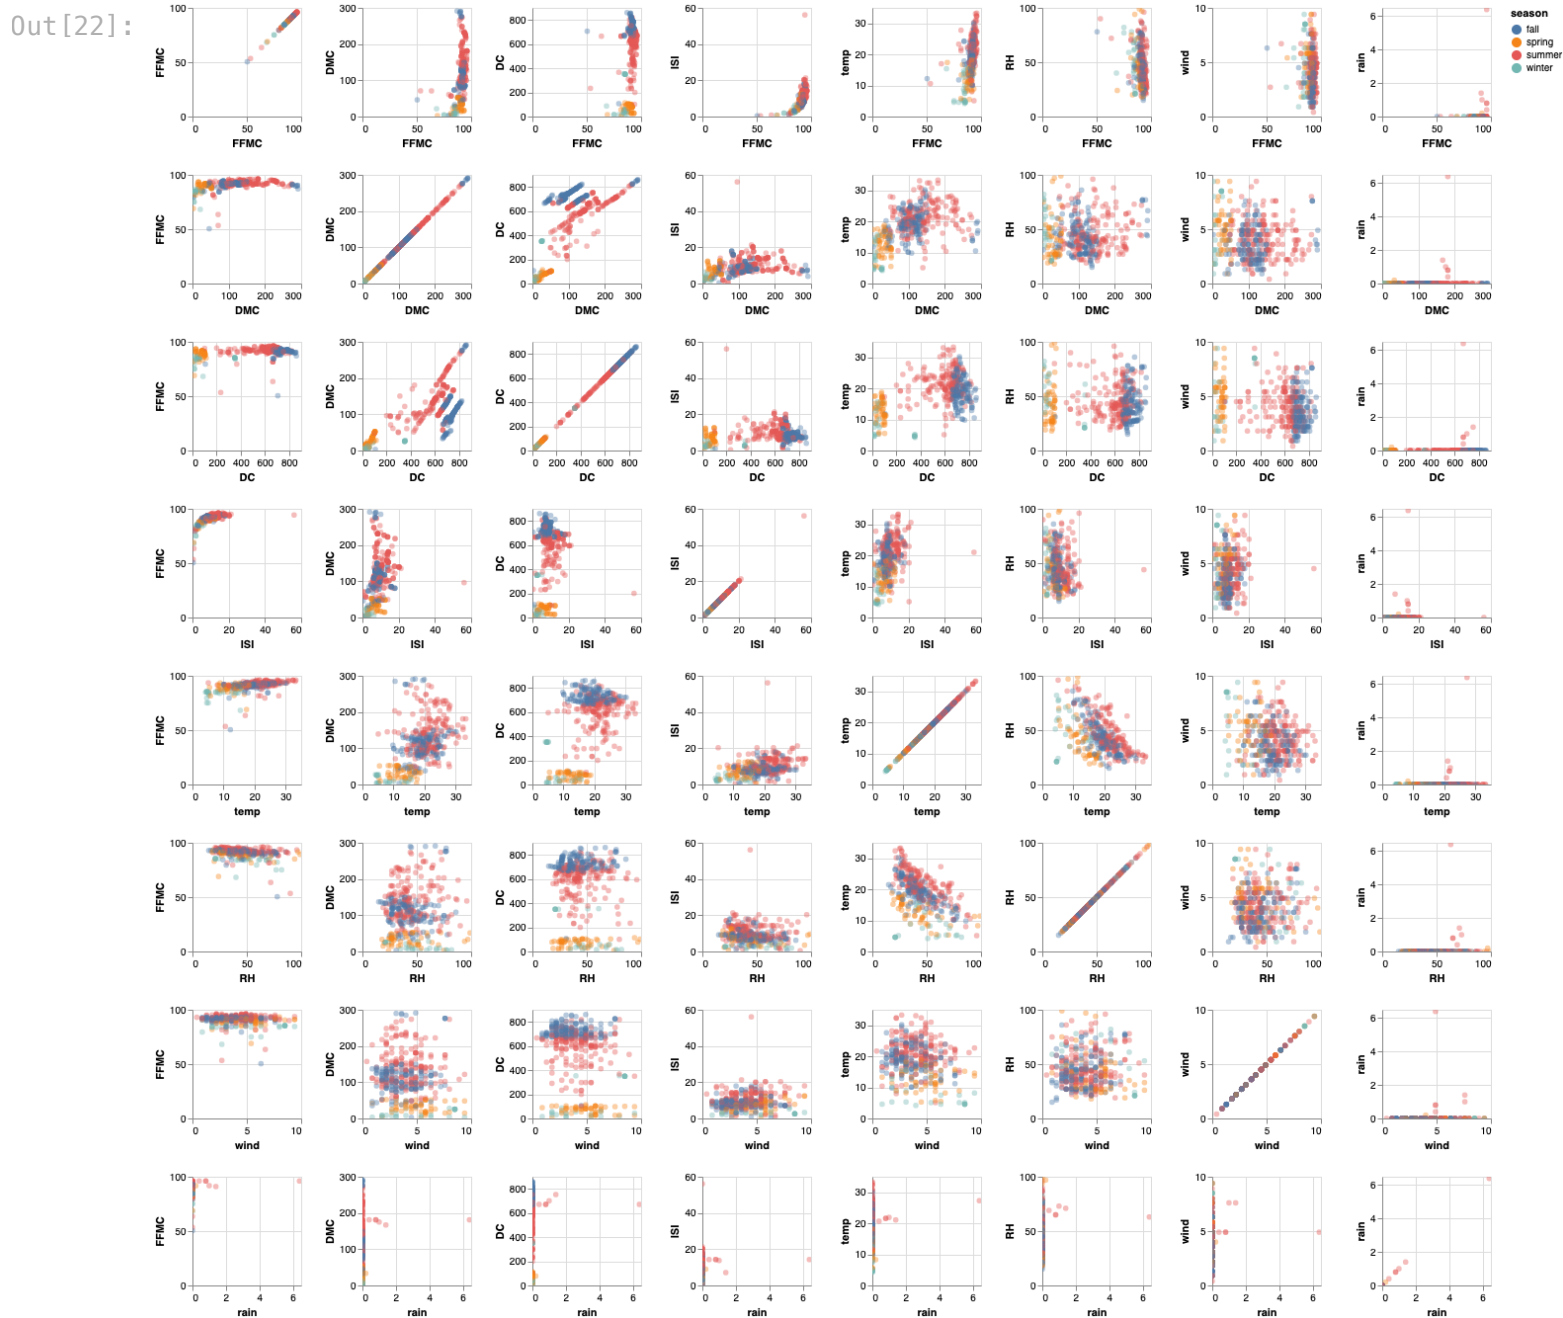


Fig 7. Pairwise relationships between numerical variables per season

Figure 7 plots the pairwise relationships between the numerical variables of the dataset. This plot shows the patterns between the numerical variables and reveals the outliers in the data. For example, the variables such as FFMC , DMC , DC , ISI and rain contain outliers. This suggests that we need to keep this in mind and deal with the outliers when making our model.

In [23]:

```

train_df_numeric = train_set.drop(["X", "Y", "month", "day"], axis=1)

corr_df = train_df_numeric.corr("spearman").stack().reset_index(name="corr")
corr_df.loc[corr_df["corr"] == 1, "corr"] = 0
corr_df["abs"] = corr_df["corr"].abs()

(
    alt.Chart(corr_df)
    .mark_circle()
    .encode(x = alt.X("level_0", title = "Variables"),
            y = alt.Y("level_1", title = "Variables"),
            size = "abs",
            color = alt.Color('corr',
                              scale = alt.Scale(scheme = 'blueorange',
                                                  domain = (-1, 1)),
                              title = "Correlation"))
    .properties(
        width = 300,
        height = 300
    )
)

```

Out[23]:




Fig 8. Correlation heatmap for numerical variables

Figure 8 shows the correlations between the numerical variables of our data. We can see that some variables seem to be correlated to each other. For example, The correlations between ISI and FFMC or the correlation between DMC and DC seem to be somewhat high. Again, we need to keep this in mind when making our model.

7. Conclusion

We have performed exploratory data analysis to gain insight about our data. We have found that our target variable is highly skewed, however after applying a log transformation, the target looks approximately normal. Therefore, we will use the log transformed target in our model. Moreover, we have found that the days of the week might not help us with the prediction of the burnt areas of the forest, and since some months had no observations, we created a seasons variable that can help us with the prediction.

Furthermore, we have detected patterns amongst our numerical predictors, however we have to be mindful of outliers and high correlations between our predictors when making our prediction model.

8. References

Dua, Dheeru, and Casey Graff. 2017. "UCI Machine Learning Repository." University of California, Irvine, School of Information; Computer Sciences. <http://archive.ics.uci.edu/ml>.

P. Cortez and A. Morais. A Data Mining Approach to Predict Forest Fires using Meteorological Data. In J. Neves, M. F. Santos and J. Machado Eds., New Trends in Artificial Intelligence, Proceedings of the 13th EPIA 2007 - Portuguese Conference on Artificial Intelligence, December, Guimaraes, Portugal, pp. 512-523, 2007. APPIA, ISBN-13 978-989-95618-0-9.